

**Feature Scaling:** Technique in which values are shifted and rescaled so that they end up ranging between 0 and 1

**Why:** Range/magnitude/units of all features in the data should be normalized so that each feature contributes approximately proportionately to the final distance. Having features on a similar scale can help the gradient descent converge more quickly towards the minima

Distance based algorithms like KNN, K-means, and SVM are most affected by the range of features. This is because behind the scenes they are using distances between data points to determine their similarity. Tree-Based Algorithms on the other hand, are fairly insensitive to the scale of the features

Most commonly used feature scaling techniques are:

1. Normalization/Min-Max Scalar
2. Standardization/ Standard scalar

What is Normalization? Normalization is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.

Here's the formula for normalization:

Normalization equation:

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

Here, Xmax and Xmin are the maximum and the minimum values of the feature respectively.

What is Standardization? Standardization is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

$$X' = \frac{X - \mu}{\sigma}$$

Here's the formula for standardization:

Note that in this case, the values are not restricted to a particular range.

Normalize or Standardize?

Normalization is good to use when you know that the distribution of your data does not follow a Gaussian distribution. This can be useful in algorithms that do not assume any distribution of the data like K-Nearest Neighbors and Neural Networks.

Standardization, on the other hand, can be helpful in cases where the data follows a Gaussian distribution. However, this does not have to be necessarily true. Also, unlike normalization,

standardization does not have a bounding range. So, even if you have outliers in your data, they will not be affected by standardization.

However, at the end of the day, the choice of using normalization or standardization will depend on your problem and the machine learning algorithm you are using. There is no hard and fast rule to tell you when to normalize or standardize your data. You can always start by fitting your model to raw, normalized and standardized data and compare the performance for best results.

It is a good practice to fit the scaler on the training data and then use it to transform the testing data. This would avoid any data leakage during the model testing process. Also, the scaling of target values is generally not required.

```
In [1]: #Implementation

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [12]: data= pd.read_csv("income.csv")
```

```
In [3]: data.shape
```

```
Out[3]: (31978, 13)
```

```
In [4]: data.head()
```

Out[4]:

	age	JobType	EdType	maritalstatus	occupation	relationship	race	gender	capitalgain	capital
--	-----	---------	--------	---------------	------------	--------------	------	--------	-------------	---------

0	45	Private	HS-grad	Divorced	Adm-clerical	Not-in-family	White	Female	0	
1	24	Federal-gov	HS-grad	Never-married	Armed-Forces	Own-child	White	Male	0	
2	44	Private	Some-college	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	
3	27	Private	9th	Never-married	Craft-repair	Other-relative	White	Male	0	
4	20	Private	Some-college	Never-married	Sales	Not-in-family	White	Male	0	

In [5]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31978 entries, 0 to 31977
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    31978 non-null  int64
1   JobType                31978 non-null  object
2   EdType                 31978 non-null  object
3   maritalstatus          31978 non-null  object
4   occupation              31978 non-null  object
5   relationship            31978 non-null  object
6   race                   31978 non-null  object
7   gender                 31978 non-null  object
8   capitalgain            31978 non-null  int64
9   capitalloss            31978 non-null  int64
10  hoursperweek           31978 non-null  int64
11  nativecountry          31978 non-null  object
12  SalStat                31978 non-null  object
dtypes: int64(4), object(9)
memory usage: 3.2+ MB
```

In [13]:

```
# filter the numeric variables from the data
data_num = data.select_dtypes(include = np.number)
data_num
```

Out[13]:

	age	capitalgain	capitalloss	hoursperweek
<b>0</b>	45	0	0	28
<b>1</b>	24	0	0	40
<b>2</b>	44	0	0	40
<b>3</b>	27	0	0	40
<b>4</b>	20	0	0	35
...	...	...	...	...
<b>31973</b>	34	594	0	60
<b>31974</b>	34	0	0	40
<b>31975</b>	23	0	0	40
<b>31976</b>	42	0	0	40
<b>31977</b>	29	0	0	40

31978 rows × 4 columns

```
In [9]: # data normalization with sklearn
from sklearn.preprocessing import MinMaxScaler

# fit scaler on training data
norm = MinMaxScaler().fit(data_num)

# transform training data
data_num_norm = norm.transform(data_num)
```

In [11]: data\_num\_norm

```
Out[11]: array([[0.38356164, 0.          , 0.          , 0.2755102 ],
 [0.09589041, 0.          , 0.          , 0.39795918],
 [0.36986301, 0.          , 0.          , 0.39795918],
 ...,
 [0.08219178, 0.          , 0.          , 0.39795918],
 [0.34246575, 0.          , 0.          , 0.39795918],
 [0.16438356, 0.          , 0.          , 0.39795918]])
```

```
In [16]: # data standardization with sklearn
from sklearn.preprocessing import StandardScaler

# fit on data column
scale = StandardScaler().fit(data_num)

# transform the data column
data_num_scale = scale.transform(data_num)

data_num_scale
```

```
Out[16]: array([[ 0.4699925 , -0.14583315, -0.21599088, -1.00589366],  
               [-1.06713224, -0.14583315, -0.21599088, -0.03384744],  
               [ 0.39679608, -0.14583315, -0.21599088, -0.03384744],  
               ...,  
               [-1.14032866, -0.14583315, -0.21599088, -0.03384744],  
               [ 0.25040325, -0.14583315, -0.21599088, -0.03384744],  
               [-0.70115016, -0.14583315, -0.21599088, -0.03384744]])
```